

# Computer Vision

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبد الله  
للعلوم والتقنية  
King Abdullah University of  
Science and Technology

KAUST Academy  
King Abdullah University of Science and Technology

February 17, 2025

1. Introduction.
2. Applications of Computer Vision.
3. Image Representation.
4. Convolutional Neural Network (CNN) and its Components.
5. CNN-based Architectures (AlexNet, VGG, InceptionNet, ResNet, EfficientNet, and MobileNet).

- ▶ Understand the basic concepts of Computer Vision and its real-world applications.
- ▶ Describe how images are represented and processed in a computer.
- ▶ Explain the fundamental building blocks of Convolutional Neural Networks (CNNs).
- ▶ Differentiate between popular CNN architectures (AlexNet, VGG, InceptionNet, ResNet, EfficientNet, and MobileNet) and their key innovations.

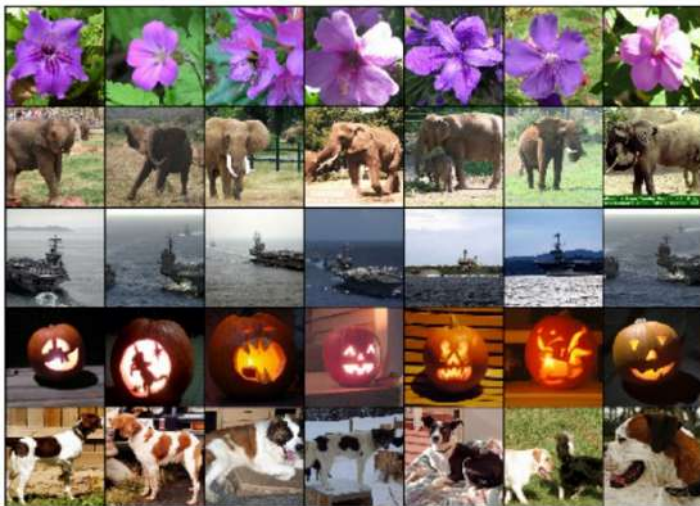
Building artificial systems that process, perceive, and reason about visual data



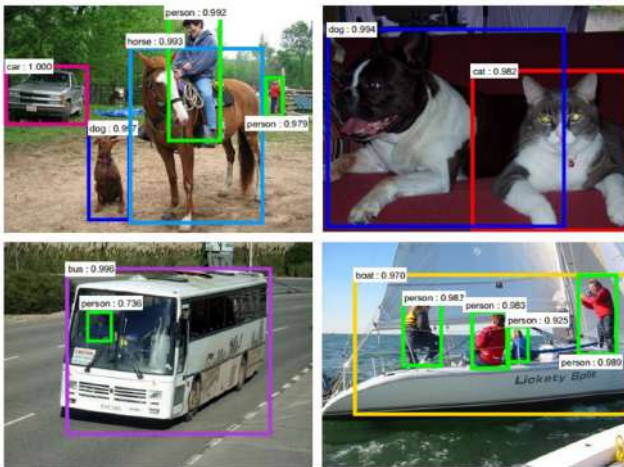
## Image Classification



## Image Retrieval




## Object Detection



Ren, He, Girshick, and Sun, 2015

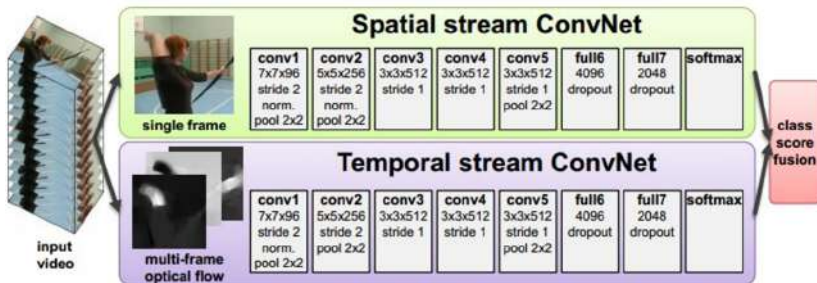


## Image Segmentation




Fabaret et al, 2012

## Video Classification



Simonyan et al, 2014


## Activity Recognition



Pose Recognition (Toshev and Szegedy, 2014)



## Medical Imaging





*A white teddy bear  
sitting in the grass*



*A man in a baseball  
uniform throwing a ball*



*A woman is holding  
a cat in her hand*

## Image Captioning

Vinyals et al, 2015

Karpathy and Fei-Fei, 2015



*A man riding a wave  
on top of a surfboard*



*A cat sitting on a  
suitcase on the floor*



*A woman standing on a  
beach holding a surfboard*

© 2015 Google LLC. All rights reserved.  
This document is the property of Google LLC. It contains confidential information and is intended only for the individual named. If you are not the named individual you should not disseminate, distribute or copy this e-mail. Please notify the sender immediately by e-mail if you have not received this e-mail by mistake. Privacy notice: [http://g.co/privacynotice](#)

## Image Generation



“Teddy bears working on new  
AI research underwater with  
1990s technology”

DALL-E 2



Style Transfer




## 3D Vision



Choy et al., 3D-R2N2: Recurrent Reconstruction Neural Network (2016)




Zhou et al., 3D Shape Generation and Completion through Point-Voxel Diffusion (2021)



Gkioxari et al., "Mesh R-CNN", ICCV 2019

# How to represent an image?

- ▶ Images are represented as Matrices with elements in  $[0, 255]$
- ▶ Grayscale images have one channel.

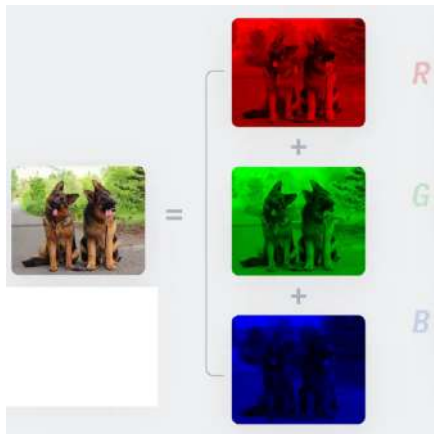


197	182	174	166	150	132	125	121	112	101	555	555
155	182	163	76	75	62	88	17	110	210	180	154
180	180	60	14	54	6	10	33	48	106	159	181
200	190	0	124	121	111	120	204	140	15	56	140
184	68	107	181	237	239	239	228	227	87	71	201
172	106	297	225	225	214	220	228	258	84	74	208
146	88	179	289	140	878	111	199	139	75	30	169
189	97	116	84	10	148	134	11	25	62	22	148
199	168	191	183	154	227	178	142	182	106	34	152
206	179	168	262	234	231	149	179	208	49	60	234
180	215	116	149	236	187	86	150	79	38	218	241
190	224	147	166	227	210	127	112	36	181	255	224
190	214	173	44	102	140	54	86	2	109	249	216
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	106	200	138	243	236
190	206	128	197	177	121	125	206	116	18	46	218

197	182	174	166	150	132	125	121	112	101	109	156
155	182	163	76	75	62	33	17	110	210	180	154
180	180	60	14	54	6	10	33	48	106	159	181
200	190	1	124	121	111	120	204	140	15	56	140
184	68	131	211	237	239	239	228	227	87	71	201
172	106	207	228	223	214	220	228	228	84	74	208
146	88	179	289	140	878	111	199	139	75	30	169
189	97	168	84	10	168	134	11	31	62	22	148
199	168	191	183	158	227	178	143	182	106	36	150
206	179	168	262	234	231	149	179	208	49	60	234
180	215	116	149	236	187	86	150	79	38	218	241
190	224	147	166	227	210	127	112	36	181	255	224
190	214	173	44	102	143	56	50	2	109	249	216
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	106	200	138	243	236
190	206	128	197	177	121	125	206	116	18	46	218

# How to represent an image?

- RGB images have 3 channels.



- So far, we've worked with tabular data, where each sample consists of structured features. We processed this data using Fully Connected Neural Networks.




Figure 12.2 Deep network architecture with multiple layers.


- But can we apply the same approach to images?

# How can we build vision models?

► Let's try.

# How can we build vision models?


- ▶ Let's consider each pixel as a feature.



- ▶ This should work! But...

# Drawbacks of Fully-Connected Neural Networks

- The number of trainable parameters becomes extremely large



# Drawbacks of Fully-Connected Neural Networks (cont.)

- ▶ Little or no invariance to shifting, scaling, and other forms of distortion
- ▶ In other words, the Fully connected NN cannot recognize that the two images (original and shifted) represent the same object (letter "A").




Figure 2: Original "A" character.




Figure 3: Shifted "A" character.



# Drawbacks of Fully-Connected Neural Networks (cont.)

- ▶ The topology of the input data is completely ignored (it treats pixels as independent features rather than structured patterns.)
- ▶ For a  $32 \times 32$  image, we have
  - Black and white patterns:  $2^{32 \times 32} = 2^{1024}$
  - Grayscale patterns:  $256^{32 \times 32} = 256^{1024}$



- ▶ We need a model that can:
  - **Find Patterns in Images:** Recognize small features like edges and shapes in different parts of the image.
  - **Work Efficiently:** Avoid looking at every pixel separately by focusing on groups of pixels together.
  - **Handle Changes:** Still recognize the same object even if it moves, rotates, or looks slightly different.

- ▶ We need a model that can:
  - **Find Patterns in Images:** Recognize small features like edges and shapes in different parts of the image.
  - **Work Efficiently:** Avoid looking at every pixel separately by focusing on groups of pixels together.
  - **Handle Changes:** Still recognize the same object even if it moves, rotates, or looks slightly different.
  - We need **Convolutional Neural Networks (CNNs)**!

# Convolutional Neural Networks (CNNs)

- ▶ CNNs are neural networks designed for image processing and consist of two main parts:
  - **Feature Extractor:** Automatically learns patterns (features) such as edges, textures, and shapes from the image.
  - **Classifier:** The extracted features are flattened into a vector and passed to a standard neural network (like the ones we used before) for classification.

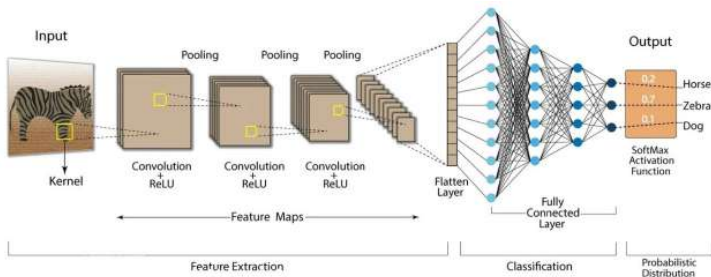


Figure 4: Convolutional Neural Network

- ▶ The feature extractor consists of three essential components:
  - **Convolution Layers:** Detect local patterns such as edges, textures, and shapes by applying filters to the image.
  - **Activation Function (ReLU):** Adds non-linearity.
  - **Pooling Layers:** Reduce the spatial size of extracted features.
- ▶ Let's start with the convolution layer.

# How Convolution Works?

- Let's consider this image.




# How Convolution Works? (cont.)

1. The image is divided into small overlapping tiles (regions).



2. We process each tile using the weights matrix of a small neural network. We call this matrix a **kernel** or **filter**.


## Processing a single tile






# How Convolution Works? (cont.)

3. Finally, the filter slides across the image (with the same weights) and processes all the tiles, creating an output **feature map**.



# How Convolution Works? (cont.)

- ▶ What we did in the last step is called the **Convolution Operation**.
- ▶ We convolved the kernel with the image, which means sliding the kernel over the entire image and computing the dot product between the kernel and small regions of the image at each step.





- This can be represented mathematically by:

$$z = W * x_{i,j} = \sum_{a=0}^{m-1} \sum_{b=0}^{n-1} W_{ab} X_{(i+a)(j+b)}$$



- $z$ : Output of the convolution at  $(i, j)$ .
- $W$ : Filter (kernel) matrix of size  $m \times n$ .
- $x_{i,j}$ : Input value at position  $(i, j)$ .
- $W_{ab}$ : Weight of the filter at  $(a, b)$ .
- $X_{(i+a)(j+b)}$ : Input value in the small region at  $(i + a, j + b)$ .

# How Convolution Works? (cont.)





The **kernel** slides across the image and produces an output value at each position

# How Convolution Works? (cont.)



The **kernel** slides across the image and produces an output value at each position


# How Convolution Works? (cont.)





The **kernel** slides across the image and produces an output value at each position

# How Convolution Works? (cont.)

- **Receptive Field:** The region of the input image that the kernel operates on at each step.
- **Feature Map:** The features extracted from the input image.



# How Convolution Works? (cont.)



We convolve multiple kernels and obtain multiple feature maps or **channels**



# How Convolution Works? (cont.)

- **Filters Example:** Different filters detect patterns like edges, textures, or smoothness, producing corresponding feature maps.

$$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$$


$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$




$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$




- ▶ Applying Convolution as such reduces the size of the borders.
- ▶ Sometimes this is not desirable.




# Controlling the Convolution Process (cont.)

- Solution: We can use **Padding**, which pads the border with zeros.
- it has two types:
  1. **Same Convolution**: Padding is added so the output size equals the input size.




2. **Full Convolution:** Padding is added so the kernel covers the entire input, including edges.


$$\text{output size} = \text{input size} + \text{kernel size} - 1$$




- **Strided Convolution:** Kernel slides along the image with a step  $> 1$
- Larger strides reduce output size.




- ▶ **Strided Convolution:** Kernel slides along the image with a step  $> 1$
- ▶ Larger strides reduce output size.



- **Pooling:** Compute mean or max over small windows to reduce resolution and extract more general features.




# Controlling the Convolution Process (cont.)





- ▶ **Dilated Convolution:** Kernel is spread out, step  $> 1$  between kernel elements.
- ▶ It expands the receptive field without increasing the number of parameters, which makes it efficient.



$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

$n_{in}$  : number of input features

$n_{out}$  : number of output features

$k$  : convolution kernel size

$p$  : convolution padding size

$s$  : convolution stride size

- ▶ Just like Fully-Connected Neural Networks, we can apply an activation over convolutional layer outputs
- ▶ It helps break linearity
- ▶ For example, Rectified Linear Unit (ReLU):  $\sigma(x) = \max(0, x)$

**Feature Map**

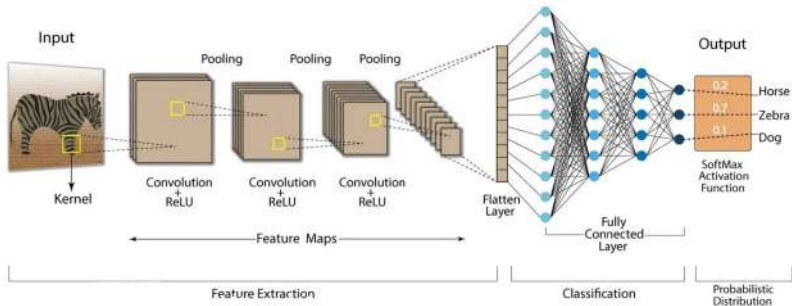
9	3	5	-8
-6	2	-3	1
1	3	4	1
3	-4	5	1

ReLU Layer




9	3	5	0
0	2	0	1
1	3	4	1
3	0	5	1


# Convolutional Neural Networks




## Convolution Layers




## Pooling Layers



## Fully-Connected Layers



## Activation Function




## Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

- ▶ Over time, researchers built advanced CNN architectures to improve performance and efficiency. These architectures introduced key innovations:
  - **AlexNet [Krizhevsky et al. 2012]**: The first CNN to achieve breakthrough performance on image classification.
  - **VGGNet [Simonyan and Zisserman, 2014]**: Used very deep networks (up to 19 layers).
  - **InceptionNet (GoogLeNet) [Szegedy et al., 2014]**: Used multiple filter sizes per layer (Inception modules).
  - **ResNet [He et al., 2015]**: Introduced skip connections for training very deep networks.
  - **EfficientNet [Tan and Le, 2019]**: Found a scaling method that simultaneously scales a CNN's depth, width, and resolution optimally using a single scaling coefficient.

- ▶ First big improvement in image classification.
- ▶ Made use of CNN, pooling, dropout, ReLU and training on GPUs.
- ▶ 5 convolutional layers, followed by max-pooling layers; with three fully connected layers at the end




- **Improvement over AlexNet:** Uses a deeper network with small filters instead of a shallow network with larger filters.
- A stack of 3x3 conv layers (vs. 11x11 conv) has same receptive field, more non-linearities, fewer parameters and deeper network representation.
- Two variants: VGG16 or VGG19 conv layers plus 3 FC layers.






- ▶ Going Deep: 22 layers
- ▶ Only 5 million parameters! ( $12\times$  fewer than AlexNet,  $27\times$  fewer than VGGNet).
- ▶ Introduced efficient "Inception module"
- ▶ Introduced "bottleneck" layers that use  $1\times 1$  convolutions to reduce the number of channels and mix information across them.




- **Inception module:** Uses multiple filter sizes ( $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$ ), in parallel, to capture different features, then combines their outputs.





- ▶ **Problem:** Making networks deeper does not always improve accuracy.
- ▶ **Why?** In very deep networks, gradients become extremely small as they move backward through layers, making learning slow or stopping it altogether (**vanishing gradient problem**).
- ▶ **Solution:** Residual Network (ResNet) introduces **skip connections (residuals)**, allowing information to flow more easily.

- ▶ Very deep networks using residual connections
- ▶ 152-layer model for ImageNet
- ▶ Stacked Residual Blocks
- ▶ **Residual:** A shortcut connection that helps the network pass information through layers more easily.





- ▶ What happens when we continue stacking deeper layers on a "plain" convolutional neural network?

- What happens when we continue stacking deeper layers on a "plain" convolutional neural network?





- What happens when we continue stacking deeper layers on a "plain" convolutional neural network?



- 56-layer model performs worse on both test and training error

- What happens when we continue stacking deeper layers on a "plain" convolutional neural network?



- 56-layer model performs worse on both test and training error
- The deeper model performs worse, but it's not caused by overfitting!





- **Fact:** Deep models have more representation power (more parameters) than shallower models.

- ▶ **Fact:** Deep models have more representation power (more parameters) than shallower models.
- ▶ **Hypothesis:** The problem is an optimization problem, deeper models are harder to optimize

- ▶ **Fact:** Deep models have more representation power (more parameters) than shallower models.
- ▶ **Hypothesis:** The problem is an optimization problem, deeper models are harder to optimize
- ▶ **Solution:** Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping

- **Fact:** Deep models have more representation power (more parameters) than shallower models.
- **Hypothesis:** The problem is an optimization problem, deeper models are harder to optimize
- **Solution:** Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping




Identity mapping:  
 $H(x) = x$  if  $F(x) = 0$

- ▶ The most extensive data for Image Classification
- ▶ 3 RGB channels from 0 to 255
- ▶ 14,197,122 images
- ▶ 1000 classes



# ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



► **Problem:** To improve accuracy, we can:

- Increase the number of layers (**depth**)
- Increase the number of neurons in each layer (**width**)
- Use higher-resolution images (**resolution**)

But finding the right balance of these three was largely based on trial and error.

► **Solution:** EfficientNet introduced a **compound scaling** method—a mathematical formula to systematically find the optimal balance among **depth**, **width**, and **resolution**.

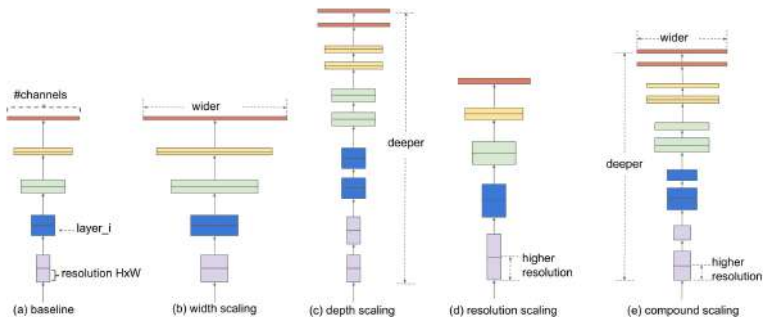


Figure 5: EfficientNet Scaling.



- ▶ **Compound Scaling** Uses a single **scaling coefficient** ( $\phi$ ) to control:
  - **Network Depth** ( $\alpha^\phi$ )  $\rightarrow$  More layers
  - **Network Width** ( $\beta^\phi$ )  $\rightarrow$  More channels per layer
  - **Input Resolution** ( $\gamma^\phi$ )  $\rightarrow$  Larger input images
- ▶ The goal: find  $\alpha, \beta, \gamma$  that balance accuracy & efficiency, then scale up optimally by increasing the global coefficient  $\phi$ .

- ▶ EfficientNet optimizes depth, width, and resolution using this constraint:

$$\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

- ▶ **Why this equation?**

- Increasing **depth** ( $\alpha$ ) increases FLOPs **linearly**.
  - Increasing **width** ( $\beta$ ) increases FLOPs **quadratically** ( $\beta^2$ ).
  - Increasing **resolution** ( $\gamma$ ) increases FLOPs **quadratically** ( $\gamma^2$ ).
- ▶ To double total FLOPs, the three factors must be balanced together.

- ▶ The authors of EfficientNet searched for the best scaling factors on a small baseline model.
- ▶ They found:

$$\alpha = 1.2, \quad \beta = 1.1, \quad \gamma = 1.15$$

- The EfficientNet family (B0 to B7) is generated using:

$$\text{Depth} = 1.2^\phi, \quad \text{Width} = 1.1^\phi, \quad \text{Resolution} = 1.15^\phi$$

Model	$\phi$	Depth ( $\alpha^\phi$ )	Width ( $\beta^\phi$ )
B0	0	$1.2^0 = 1.0$	$1.1^0 = 1.0$
B1	1	$1.2^1 = 1.2$	$1.1^1 = 1.1$
B2	2	$1.2^2 = 1.44$	$1.1^2 = 1.21$
B3	3	$1.2^3 = 1.73$	$1.1^3 = 1.33$
B4	4	$1.2^4 = 2.07$	$1.1^4 = 1.46$
B5	5	$1.2^5 = 2.49$	$1.1^5 = 1.61$
B6	6	$1.2^6 = 2.99$	$1.1^6 = 1.77$
B7	7	$1.2^7 = 3.58$	$1.1^7 = 1.94$

Table 1: Scaling EfficientNet from B0 to B7

- We multiply these scaling factors by the baseline EfficientNet-B0 values and round to the nearest integer to get the new depth, width, and resolution for each model.

- ▶ EfficientNet models achieve state-of-the-art accuracy with significantly fewer parameters and FLOPs.

- ▶ EfficientNet models achieve state-of-the-art accuracy with significantly fewer parameters and FLOPs.
- ▶ EfficientNet-B7 reaches 84.4% Top-1 and 97.3% Top-5 accuracy on ImageNet.

- ▶ EfficientNet models achieve state-of-the-art accuracy with significantly fewer parameters and FLOPs.
- ▶ EfficientNet-B7 reaches 84.4% Top-1 and 97.3% Top-5 accuracy on ImageNet.
- ▶ More efficient than previous CNN models—8.4x smaller and 6.1x faster than competitors.




Figure 6: EfficientNet Performance on Imagenet.



## ► Why MobileNets?

- Small-sized models are crucial for mobile and embedded devices.
- MobileNets reduce computational cost and memory usage while maintaining good accuracy.

## ► Key Idea:

- Use **depthwise-separable convolutions** to significantly reduce computation compared to standard convolutions.

- **Computational cost of standard convolution:**

$$\text{Cost} = \# \text{ filter params} \times \# \text{ filter positions} \times \# \text{ filters}$$

- Filters operate on all input channels, increasing computation significantly.




- ▶ Split standard convolution into two steps:
  - **Depthwise Convolution:** Applies a single filter per input channel.
  - **Pointwise Convolution:** Combines outputs from depthwise convolution.
- ▶ **Key Benefit:** Reduces computational cost significantly compared to standard convolution.


Normal Convolution



Depthwise Separable Convolution



- Operates on each input channel separately.



- Combines outputs from depthwise convolution using  $1 \times 1$  convolutions (mixes channels).




$$\text{Computational cost} = \# \text{filter params} \times \# \text{filter positions} \times \# \text{of filters}$$


## ► MobileNet v2:

- Adds **residual connections**.
- Introduces:
  - **Expansion step:** Expands input dimensions before depthwise convolution.
  - **Projection step:** Reduces dimensions after processing.

MobileNet v1



MobileNet v2



These slides have been adapted from

- ▶ Fei-Fei Li, Yunzhu Li & Ruohan Gao, Stanford CS231n: [Deep Learning for Computer Vision](#)
- ▶ Assaf Shocher, Shai Bagon, Meirav Galun & Tali Dekel, WAIC DL4CV [Deep Learning for Computer Vision: Fundamentals and Applications](#)
- ▶ Justin Johnson, UMich EECS 498.008/598.008: [Deep Learning for Computer Vision](#)
- ▶ Sander Dieleman, Deepmind: [Deep Learning Lecture Series 2020](#)